

# A Lightweight, Swappable Voxel Ray-Tracer

Ann Gao, Nathan Hall, Jonah Henriksson, Christopher Jenkins, Felipe Schmidt

**Abstract:** Voxel rendering is a technique for visualizing 3D volumes, with applications in medical technology, terrain visualization, and video games. Simple voxel rendering implementations rely on “meshing”, a process by which the volume is converted to a mesh for rasterization. For many tasks, this is sufficient, but also suffers from memory constraints which prevents meshing from scaling to large scenes (since both the voxel data and heavy mesh representation has to be stored). An alternative solution that lacks such constraints is ray-tracing; which operates on the voxel data itself. In this project, we present a CPU voxel ray-tracer implemented in Rust, alongside two storage solutions that can be swapped: a dense array-backed storage, and a sparse octree storage. <Talk about measuring results between the two>

Terms: Voxels, Ray-Tracing, Octree

## Introduction

The term “voxels” is derived from the term “pixels”: where a “pixel” is a “**p**icture **e**lement” that represents part of a picture, a “voxel” is a “**v**olume **e**lement” that represents part of a volume. Volumetric data can range from MRI slices to point clouds to generated volumes in video games. Voxels are a representation of this volumetric data in a grid structure, just as pixels store raster images in a grid structure. However, while pixels and voxels are similar in structure, the way they are visualized differs vastly. Pixels can simply be drawn to a screen surface, but voxels are points in 3D space and drawing techniques can either draw the individual points, as with point cloud rendering, or the isosurface, which is the concern of this paper.

Isosurface rendering is the ubiquitous form of voxel rendering, to the point that “voxel rendering” almost always refers to rendering the isosurface of a voxel grid. The isosurface of a voxel grid is the surface present between voxels that surpass some threshold and those which don’t (i.e. the boundary between voxels that are “present” and those which aren’t). To render such an isosurface, the approach varies on the underlying render technique used. Such render techniques generally fall into two categories: rasterization and ray-tracing.

Ray-tracing is the earliest method of rendering, which works by simulating rays of light in a scene, making it capable of rendering realistic images. However, it is computationally expensive, leading to the popularity of another method for real time applications: rasterization. Rasterization works by projecting triangles of a polygonal mesh onto the camera’s 2D view and filling in their bounds using small programs called “shaders”. Because of the strengths and weaknesses of both methods, voxel rendering techniques fall into two categories as well:

isosurface mesh extraction for rasterization and voxel traversal for ray-tracing (there are also hybrid techniques, but they exist as an optimization of voxel traversal).

Isosurface mesh extraction, popularly known as “meshing”, is the process of converting voxels to a mesh of the isosurface, and there are many meshing algorithms, such as “greedy meshing”, “marching cubes”, “naive surface nets”, etc. All have their own characteristics making them suitable for particular applications. However, when rendering large scenes, meshing and rasterization typically struggle to scale in performance. Meanwhile, ray tracing is capable of scaling in performance, due to optimizations that are relatively easy to implement with ray tracing, such as acceleration structures and LOD (level-of-detail), and the lack of a mesh generation step when updating voxels, as well as better memory scaling (a single voxel may be represented as a single byte, compared to a single naive block mesh which has 8 vertices, each represented by 3 32-bit floating point numbers, on top of the underlying voxel data). Because of this difference in scaling potential, when rendering large scenes on mid to high end consumer computers, ray tracing becomes more performant than rasterization.

The goal of this project is to implement a CPU voxel ray-tracer in the Rust language with two storage backends: a dense array storage and a sparse voxel octree storage. The dense array represents a naive voxel storage implementation, while the sparse octree represents an acceleration structure implementation: an optimization for ray traversal via a space partitioning structure. Using both storage solutions, we can assess the benefits of octrees for optimizing ray traversal in voxel scenes.

## Methodology

The ray tracer is divided into the following parts:

- A voxel generator for constructing scenes
- The storage backend for scenes (either dense or sparse)
- The ray tracer that renders a scene
- The framebuffer that the ray tracer writes to and saves an image

<Elaborate more on how they tie into each other?>

To swap between the two storage backends, we have adopted an object oriented design: a scene interface will be implemented for both solutions which allows for instantiating the storage object from a voxel generator and tracing a ray through the scene to get a voxel.

In order to benchmark and compare both storage backends, we are using the “criterion” crate, which provides a harness we can use to measure the time to render via both backends. To further assess the performance, we also are using the “tracing” crate, which is used to measure the performance of individual functions and produce flame graphs of the application.

To verify the results, the framebuffer generated by rendering is capable of being written to an image file using the “image” crate.

## Implementation

<Elaborate on ray-tracing>

<Camera setup>

<Ray instantiation>

<Ray task>

<Elaborate on voxel representation and generation>

<Voxel generation source>

<Elaborate on storage solutions>

The storage backends are constructed using a voxel generator and an axis-aligned bounding box (AABB). An AABB describes the position and extents of a volume. By iterating over every coordinate in the AABB, the storage backend can query the voxel generator for a voxel and store it if present. The AABB can also be used during a ray trace, as a quick bounds check.

<Dense: Fast Traversal Algorithm>

For the dense storage backend, there exist many algorithms for ray tracing through a multidimensional array, but most are based on the “Fast Voxel Traversal” algorithm by Amanatides and Woo, which will be what we use for ray tracing in the dense storage backend. This algorithm works by tracking the index and position for each axis and moving the ray along the axis with the shortest distance to the next index. By doing so, it ensures that no voxels are missed during iteration. When a voxel is found or the indices go out of bounds, it returns.

<Sparse: Octree space partitioning>

In the world of 2D graphics there is a data structure known as a quadtree which is used to divide a 2D plane into 4 quadrants one to represent top left, top right, bottom left, and bottom right. These quadrants are stored as children of a parent node and so on and so forth. An Octree is the 3D equivalent of this. Within each Octree is a root that represents an area in 3D space, within this root there are 8 children which are 8 subdivisions of the parent root, each subsequent child breaks down the space until the children being stored are individual voxels.

An octree is a helpful tool in rendering 3D spaces as they are a very compact way of storing voxels without wasting memory. They are also very quick to traverse as one can navigate to a point in a 3D space by either descending or ascending the structure to pinpoint any given location.

One of the key features is the PreOrder Traversal, this is helpful because it can help to determine line of sight for an object, or for our purposes, a ray. Since the ray can check what objects it will or will not collide with by simply traversing the octree it greatly reduces a lot of redundant calculations that would have been performed were the voxels to be stored in a different data structure.

<Elaborate on framebuffer>

The framebuffer is a data structure that holds all the image information of the rendered scene. It contains a memory reference to a heap-allocated array of color values for each pixel. It allows for multiple threads to access different pixels concurrently to allow faster image generation.

## Evaluation

<Show results for both storage solutions>

<Discuss results>

## Conclusion